

QTM 220 Final Exam (12/7)

Problem 1

Suppose you have a sample $Y_1, Y_2, \dots, Y_n$ drawn with replacement from a population with mean μ and standard deviation σ .

Consider the following estimator for μ : $\tilde{Y} = \frac{1}{2n} \sum_{i=1}^n Y_i$

Part A

What is $E[\tilde{Y}]$?

$$E[\tilde{Y}] = E\left[\frac{1}{2n} \sum_{i=1}^n Y_i\right]$$

$$= \frac{1}{2n} E\left[\sum_{i=1}^n Y_i\right]$$

$$= \frac{1}{2n} \sum_{i=1}^n E[Y_i]$$

by linearity of expectations

$$= \frac{1}{2n} \sum_{i=1}^n E[Y]$$

by IID

$$= \frac{1}{2n} \sum_{i=1}^n \mu$$

by definition

$$= \frac{1}{2n} n \mu$$

$$= \frac{n}{2n} \mu$$

$$= \frac{1}{2} \mu$$

Part B

How does $E[\tilde{Y}]$ compare to $E[\bar{Y}]$? (where $\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$)

similar steps show that $E[\bar{Y}] = \mu$. $\bar{Y}$ is unbiased, $\tilde{Y}$ is biased; bias does not change in n

Part C

Showing your work, calculate $Var[\tilde{Y}]$.

$$Var[\tilde{Y}] = V[\frac{1}{2n} \sum_{i=1}^n Y_i]$$

$$= (\frac{1}{2n})^2 V[\sum_{i=1}^n Y_i]$$

$$= (\frac{1}{2n})^2 \sum_{i=1}^n V[Y]$$

because IID, covariances are 0

$$= (\frac{1}{2n})^2 \sum_{i=1}^n \sigma^2$$

by definition

$$= (\frac{1}{2n})^2 n \sigma^2$$

$$= \frac{n}{(2n)^2} \sigma^2$$

$$= \frac{1}{4n} \sigma^2$$

Part D

How does $Var[\tilde{Y}]$ compare to $Var[\bar{Y}]$?

Similar steps show that $V[\tilde{Y}] = \frac{\sigma^2}{n}$

$\frac{1}{n} > \frac{1}{4n}$ so our $\tilde{Y}$ estimator has a lower variance than the sample mean. But the estimator is biased, so we probably wouldn't want to use it over the mean.

Problem 2

Suppose we have a sample of $Y_1 \dots Y_n$ and binary $X_1 \dots X_n$ stored in an **R** dataframe 'df' with n rows. We run this code.

```
B=10000
theta<- rep(NA,B)
for( b in 1:B ){
  bootind = sample(1:n,size=n,replace=T)
  dfboot = df[bootind,]
  theta[b]<- (mean(dfboot$y[dfboot$x==1])-mean(dfboot$y[dfboot$x==0]))
}
```





Figure 1: Histogram of the vector Y 's elements Figure 2: Jittered Scatterplot of X and Y




Figure 3: Histogram of the theta's elements

Part a

What does this code calculate? State it in words and mark it with an 'a', and an arrow or bracket if necessary for clarity, on the appropriate plot above.

```
mean(Y)
```

The sample mean. It's indicated by the solid vertical line on the sample histogram.

Part b

What does this code calculate? State it in words and mark it with a 'b', and an arrow or bracket if necessary for clarity, on the appropriate plot above.

```
sd(Y)
```

The sample standard deviation of Y. It's the distance between the solid vertical and either dashed vertical line on the sample histogram.

Part c

What does this code calculate? State it in words and mark it with a 'c', and an arrow or bracket if necessary for clarity, on the appropriate plot above.

```
mean(Y[X==0])
```

The conditional sample mean of Y among X=0. It is the lower dashed line on the scatterplot.

Part d

What does this code calculate? State it in words and mark it with a ‘d’, and an arrow or bracket if necessary for clarity, on the appropriate plot above.

```
sd(theta)
```

The bootstrap sampling distribution’s standard deviation for the difference in the mean of y between x=1 and x=0. Estimated standard error. It’s the distance between the solid vertical and either dashed vertical line on the theta sampling distribution histogram.

Part e

What does this code calculate? State it in words and mark it with a ‘e’, and an arrow or bracket if necessary for clarity, on the appropriate plot above.

```
(mean(y[x==1])-mean(y[x==0])) + 1.96*(sd(y[x==1])/sqrt(n1) + sd(y[x==0])/sqrt(n0))  
(mean(y[x==1])-mean(y[x==0])) - 1.96*(sd(y[x==1])/sqrt(n1) + sd(y[x==0])/sqrt(n0))
```

A 95% confidence interval — based on normal approximation — for the difference in means for the population from which this sample was drawn. This is indicated by the horizontal segment (with a dot in the middle) on the sampling distribution histogram (theta).

Part f

How could you calculate roughly the same thing as in part e using ‘sd(theta)’

We can use the standard deviation of the bootstrap sampling distribution as our estimate of standard error.

```
(mean(y[x==1])-mean(y[x==0])) + 1.96 * sd(theta)  
(mean(y[x==1])-mean(y[x==0])) - 1.96 * sd(theta)
```

conceptual addition: why no n? What is standard error? talk through bootstrap in words?

Problem 3

change to quadratic!

Problem 4

We run an experiment in which a treatment, D , is randomly assigned to participants conditional on sex. The table below shows our data, including our outcome of interest, Y .

Sex	Treatment (D)	Outcome (Y)	Y(1)	Y(0)
Male	1	20		
Male	0	5		
Male	0	10		
Male	0	5		
Female	1	10		
Female	1	5		
Female	0	5		
Female	0	0		
Female	0	5		
Female	1	10		
Female	0	5		
Female	1	5		

Part A

Fill in the table with potential outcomes based on the observed data. Put a question mark (?) in the cells for which you do not have data.

Part B

What is the probability of receiving treatment for males? What is the probability of receiving treatment for females?

$$P(D = 1|Male) = \frac{1}{4}$$

$$P(D = 1|Female) = \frac{1}{2}$$

Part C

What is the treatment effect among males? What is the treatment effect among females?

$$\text{Males: } 20 - \frac{5+10+5}{3} = 20 - \frac{20}{3} = 13.333$$

$$\text{Females: } \frac{10+5+10+5}{4} - \frac{5+0+5+5}{4} = \frac{30}{4} - \frac{15}{4} = 7.5 - 3.75 = 3.75$$

Part D

If you ignored the conditional randomization and calculated the average treatment effect without accounting for sex, what would your estimate be?

$$\frac{20+10+5+10+5}{5} - \frac{5+10+5+5+0+5+5}{7} = \frac{50}{5} - \frac{35}{7} = 10 - 5 = 5$$

Part E

To analyze our data, we are interested in fitting the following model:

$$Y_i = \beta_0 + \beta_1 D_i + \beta_2 Sex_i + \beta_3 D_i : Sex_i$$


where $Sex_i = 1$ for males and $= 0$ for females. Fill in the table below using β s

	β s
Avg Y [Treated, Male]	$\beta_0 + \beta_1 + \beta_2 + \beta_3$
Avg Y [Control, Male]	$\beta_0 + \beta_2$
Avg Y [Treated, Female]	$\beta_0 + \beta_1$
Avg Y [Control, Female]	β_0
$\hat{ATE}_{Male}$	$\beta_1 + \beta_3$
$\hat{ATE}_{Female}$	β_1

conceptual addition: when we run this regression in R, why might we be concerned about using the standard errors `lm()` gives us? [homoskedasticity assumption]

Problem 5

Consider the following data with four observations.



Our goal is to use leave one out cross-validation to choose between a linear regression model and a quadratic regression model for this data. Below we have plotted these data four times. In the top left panel, the left-most observation has been left out (denoted by an open circle). In the top right, the second observation has been left out. In the bottom left, the third observation has been left out. In the bottom right, the right-most observation has been left out.

Part A





On the figure above, draw a linear model for each of the plots based on the non-left-out data.

Part B

On the figure above, draw the residuals from the linear models to the left-out observation for each of the plots.

Part C

On the figure above, draw a quadratic regression line (i.e., line that can be parabolic) for each of the plots based on the non-left-out data.



Part D

On the figure above, draw the residuals from the quadratic models to the left-out observation for each of the three plots.

Part E

Based on the residuals you calculated in parts B and D, which model is preferred?

Problem 6

piece-wise constant

Problem 7

omitted variable bias with interaction

Part A

Suppose we want β_1 from $y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 x_i : z_i + e_i$ but we run the regression $lm(y \sim x)$ without z or the interaction term. We have de-meaned our variables x_i and z_i such that $\bar{x} = \sum_{i=1}^n x_i = 0$ and $\bar{z} = \sum_{i=1}^n z_i = 0$. Derive what your coefficient on x from $lm(y \sim x)$, which is $\frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2}$, will be in terms of β s and justify your steps. Hint: you can relabel $x_i : z_i$ as a different variable.

The coefficient on x from $lm(y \sim x)$ is

$$\frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2}$$

Plug in y_i defined by the multivariate regression: $y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 x_i : z_i + e_i$

Let's redefine $x_i : z_i = w_i$

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 w_i + e_i$$

$$\frac{\sum_{i=1}^n x_i (\beta_0 + \beta_1 x_i + \beta_2 z_i + \beta_3 w_i + e_i)}{\sum_{i=1}^n x_i^2}$$

Distribute

$$\frac{\sum_{i=1}^n x_i \beta_0}{\sum_{i=1}^n x_i^2} + \frac{\sum_{i=1}^n x_i \beta_1 x_i}{\sum_{i=1}^n x_i^2} + \frac{\sum_{i=1}^n x_i \beta_2 z_i}{\sum_{i=1}^n x_i^2} + \frac{\sum_{i=1}^n x_i \beta_3 w_i}{\sum_{i=1}^n x_i^2} + \frac{\sum_{i=1}^n x_i e_i}{\sum_{i=1}^n x_i^2}$$

Let's work through each term.

$$\frac{\sum_{i=1}^n x_i \beta_0}{\sum_{i=1}^n x_i^2} = \frac{\beta_0 \sum_{i=1}^n x_i}{\sum_{i=1}^n x_i^2} = \frac{\beta_0 * 0}{\sum_{i=1}^n x_i^2} = 0$$

as x is demeaned, $\sum x_i = 0$

$$\frac{\sum_{i=1}^n x_i \beta_1 x_i}{\sum_{i=1}^n x_i^2} = \frac{\beta_1 \sum_{i=1}^n x_i x_i}{\sum_{i=1}^n x_i^2} = \beta_1$$

$$\frac{\sum_{i=1}^n x_i \beta_2 z_i}{\sum_{i=1}^n x_i^2} = \frac{\beta_2 \sum_{i=1}^n x_i z_i}{\sum_{i=1}^n x_i^2}$$

$$\frac{\sum_{i=1}^n x_i \beta_3 w_i}{\sum_{i=1}^n x_i^2} = \frac{\beta_3 \sum_{i=1}^n x_i w_i}{\sum_{i=1}^n x_i^2}$$

$$\frac{\sum_{i=1}^n x_i e_i}{\sum_{i=1}^n x_i^2} = \frac{0}{\sum_{i=1}^n x_i^2} = 0$$

we know that $\sum_{i=1}^n x_i e_i = 0$ from the first order conditions on our least squares estimator (we chose the intercept and coefficient(s) that minimize the sum of squared residuals)

Thus we have

$$\frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2} = \beta_1 + \frac{\beta_2 \sum_{i=1}^n x_i z_i}{\sum_{i=1}^n x_i^2} + \frac{\beta_3 \sum_{i=1}^n x_i w_i}{\sum_{i=1}^n x_i^2}$$

Plugging back in for w_i ,

$$\frac{\sum_{i=1}^n x_i y_i}{\sum_{i=1}^n x_i^2} = \beta_1 + \frac{\beta_2 \sum_{i=1}^n x_i z_i}{\sum_{i=1}^n x_i^2} + \frac{\beta_3 \sum_{i=1}^n x_i z_i x_i}{\sum_{i=1}^n x_i^2}$$

conceptual question: what happens if x is randomized?