

Lecture 1

Descriptive Regression with a Binary Covariate

Outline

- ▶ Univariate Descriptive Statistics
- ▶ Descriptive Regression with a Binary Covariate

What are descriptive statistics?

- ▶ Descriptive statistics are summaries of the data we have.

What are descriptive statistics?

- ▶ Descriptive statistics are summaries of the data we have.
- ▶ The data we have might be a sample (e.g., an exit poll) or it might be the whole population (e.g., administrative vote counts), sometimes known as a census.

What are descriptive statistics?

- ▶ Descriptive statistics are summaries of the data we have.
- ▶ The data we have might be a sample (e.g., an exit poll) or it might be the whole population (e.g., administrative vote counts), sometimes known as a census.
- ▶ Having the whole population is becoming more common in the era of Big Data.

What are descriptive statistics?

- ▶ Descriptive statistics are summaries of the data we have.
- ▶ The data we have might be a sample (e.g., an exit poll) or it might be the whole population (e.g., administrative vote counts), sometimes known as a census.
- ▶ Having the whole population is becoming more common in the era of Big Data.
- ▶ Note: if you are interested in the causal effects of a treatment (a drug, a policy, etc.), you can never really have the whole population (we will discuss this in a few weeks).

What makes a good summary?

What are the two things that make a good summary?

What makes a good summary?

What are the two things that make a good summary?

From the google definition:

1. brief
2. includes main points

What makes a good descriptive statistic?

1. brief: at most a few numbers

What makes a good descriptive statistic?

1. brief: at most a few numbers
2. includes main points:

What makes a good descriptive statistic?

1. brief: at most a few numbers
2. includes main points:
 - ▶ sufficient to describe main characteristics of the data
 - ▶ there is a technical definition of sufficiency that I can discuss in office hours for those interested

What makes a good descriptive statistic?

1. brief: at most a few numbers
2. includes main points:
 - ▶ sufficient to describe main characteristics of the data
 - ▶ there is a technical definition of sufficiency that I can discuss in office hours for those interested
 - ▶ minimizes the information lost by reducing the data
 - ▶ I'm using "information" in a loose sense here, will formalize a bit later
 - ▶ reducing the data means taking the data, which are often represented by $n * v$ numbers (with n the number of observations and v the number of variables) and reducing to a descriptive statistic that is represented by a few numbers.

What makes a good descriptive statistic?

1. brief: at most a few numbers
2. includes main points:
 - ▶ sufficient to describe main characteristics of the data
 - ▶ there is a technical definition of sufficiency that I can discuss in office hours for those interested
 - ▶ minimizes the information lost by reducing the data
 - ▶ I'm using "information" in a loose sense here, will formalize a bit later
 - ▶ reducing the data means taking the data, which are often represented by $n * v$ numbers (with n the number of observations and v the number of variables) and reducing to a descriptive statistic that is represented by a few numbers.

Enough forest. Trees (examples) are coming!



"Hold on, where's the forest again?"

It is easy in a statistics class to focus on the specifics. I want you to remember why we are doing what we are doing.

Outline

Univariate Descriptive Statistics

Univariate Example: SCOTUS data

- ▶ Data on 27 justices from the Warren ('53 - '69), Burger ('69 - '86), and Rehnquist ('86 - '05) courts

Univariate Example: SCOTUS data


- ▶ Data on 27 justices from the Warren ('53 - '69), Burger ('69 - '86), and Rehnquist ('86 - '05) courts
- ▶ Data can be interpreted as a census for the 1953 - 2005 era.

Univariate Example: SCOTUS data

- ▶ Data on 27 justices from the Warren ('53 - '69), Burger ('69 - '86), and Rehnquist ('86 - '05) courts
- ▶ Data can be interpreted as a census for the 1953 - 2005 era.
- ▶ One Variable: Percentage of votes in liberal direction for each justice in civil liberties cases

Justice	CLlib
Black	74.2
Reed	33.3
Thomas	25.6

Describing Univariate Data



Description/Summarization

- ▶ Even for univariate data, we may want to focus on one or two characteristics of the population instead of describing it in its entirety.


Description/Summarization

- ▶ Even for univariate data, we may want to focus on one or two characteristics of the population instead of describing it in its entirety.
- ▶ We typically focus on measures of center and spread.

Description/Summarization


- ▶ Even for univariate data, we may want to focus on one or two characteristics of the population instead of describing it in its entirety.
- ▶ We typically focus on measures of center and spread.
- ▶ We will discuss the rationale for this over the coming weeks.

Univariate Activity



Find the rough location of the mean and median.


Univariate Activity Solution



Mean/Median Discussion

In pairs, discuss what the mean and median represent in the context of the CLib variable for these justices. What does it mean for the mean to be larger than the median in this context?

Mean as a balance point



Choosing a Descriptive Statistic

How might we choose between these descriptive statistics?

1. brief: at most a few numbers

Choosing a Descriptive Statistic

How might we choose between these descriptive statistics?

1. brief: at most a few numbers
2. includes main points:

Choosing a Descriptive Statistic

How might we choose between these descriptive statistics?

1. brief: at most a few numbers
2. includes main points:
 - ▶ sufficient to describe main characteristics of the data
 - ▶ minimizes the information lost by reducing the data

A Least Squares Summary

Suppose we want to pick a single number (m) for the middle of the data that summarizes all the values for y , by minimizing the sum of squared residuals (i.e., least squares).

$$SSR(\tilde{m}) = \sum_{i=1}^N (y_i - \tilde{m})^2$$

A Least Squares Summary


Suppose we want to pick a single number (m) for the middle of the data that summarizes all the values for y , by minimizing the sum of squared residuals (i.e., least squares).

$$SSR(\tilde{m}) = \sum_{i=1}^N (y_i - \tilde{m})^2$$

Activity: Using calculus, derive the least squares estimator

1. Calculate the derivative of SSR with respect to $\tilde{m}$
2. Set the derivative equal to 0
3. Solve for m

The Objective Function for SSR CLlib




1.


$$\begin{aligned} SSR(\tilde{m}) &= \sum_{i=1}^N (y_i - \tilde{m})^2 \\ &= \sum_{i=1}^N (y_i^2 - 2y_i\tilde{m} + \tilde{m}^2) \end{aligned}$$

$$\frac{\partial SSR(\tilde{m})}{\partial \tilde{m}} = \sum_{i=1}^N (-2y_i + 2\tilde{m})$$


The Slope of the Tangent Line for SSR CLlib



The Slope of the Tangent Line for SSR CLlib



The Slope of the Tangent Line for SSR CLlib



1.

$$\begin{aligned} SSR(\tilde{m}) &= \sum_{i=1}^N (y_i - \tilde{m})^2 \\ &= \sum_{i=1}^N (y_i^2 - 2y_i\tilde{m} + \tilde{m}^2) \\ \frac{\partial S(m)}{\partial \tilde{m}} &= \sum_{i=1}^N (-2y_i + 2\tilde{m}) \end{aligned}$$

2.

$$0 = \sum_{i=1}^N (-2y_i + 2m)$$

3.

$$m = \frac{1}{N} \sum_{i=1}^N y_i$$

Therefore, the mean is a least squares summary.

An Least Absolute Value Summary

Suppose we want to pick a single number (m) for the middle of the data that summarizes all the values for y , by minimizing the sum of absolute deviations.

$$SAV(\tilde{m}) = \sum_{i=1}^N |y_i - \tilde{m}|$$

An Least Absolute Value Summary

Suppose we want to pick a single number (m) for the middle of the data that summarizes all the values for y , by minimizing the sum of absolute deviations.

$$SAV(\tilde{m}) = \sum_{i=1}^N |y_i - \tilde{m}|$$

- ▶ Why is this harder to solve?
- ▶ Any guesses?

Why we focus on means

- ▶ Means often provide a “good” summary of the data (and it is relatively easy to tell when they provide bad summaries).

Why we focus on means

- ▶ Means often provide a “good” summary of the data (and it is relatively easy to tell when they provide bad summaries).
- ▶ The properties of means are relatively easy to describe.

Why we focus on means

- ▶ Means often provide a “good” summary of the data (and it is relatively easy to tell when they provide bad summaries).
- ▶ The properties of means are relatively easy to describe.
- ▶ Randomized experiments only identify mean causal effects (as do regression and matching when the appropriate assumptions hold).

Variance and Standard Deviation

- ▶ Mean is the least squares predictor in terms of SSR.
- ▶ How small is “least” (i.e., how much spread around the mean)?
- ▶ Variance and Standard Deviation are useful transformations of SSR for the mean.

Variance and Standard Deviation

- ▶ Mean is the least squares predictor in terms of SSR.
- ▶ How small is “least” (i.e., how much spread around the mean)?
- ▶ Variance and Standard Deviation are useful transformations of SSR for the mean.

Variance

$$\begin{aligned} S^2 &= \frac{\sum_{i=1}^N (y_i - \bar{y})^2}{N - 1} \\ &= \frac{SSR}{N - 1} \end{aligned}$$

Variance and Standard Deviation

- ▶ Mean is the least squares predictor in terms of SSR.
- ▶ How small is “least” (i.e., how much spread around the mean)?
- ▶ Variance and Standard Deviation are useful transformations of SSR for the mean.


Variance

$$\begin{aligned} S^2 &= \frac{\sum_{i=1}^N (y_i - \bar{y})^2}{N - 1} \\ &= \frac{SSR}{N - 1} \end{aligned}$$


Standard Deviation

$$S = \sqrt{S^2}$$

Mean and Standard Deviation



Variance as turning force



Why Variance and Standard Deviation?

- ▶ Natural measures of accuracy for means.


Why Variance and Standard Deviation?

- ▶ Natural measures of accuracy for means.
- ▶ Means and standard deviations are sufficient to describe the distribution of normally distributed data.


Why Variance and Standard Deviation?

- ▶ Natural measures of accuracy for means.
- ▶ Means and standard deviations are sufficient to describe the distribution of normally distributed data.
- ▶ For non-normal data, means and standard deviations provide information via Chebyshev's inequality and other concentration inequalities.

Normally-Shaped Histograms



Chebyshev for Non-Normal Histograms



Special Case: Binary Variables

Justice	Underrepresented Group
Black	0
Marshall	1
Burger	0
O'Connor	1
Scalia	0
Thomas	1
Proportion	3/27

Mean and Variance for UR Group Variable

$$\begin{aligned}\bar{y} &= \frac{1}{N} \sum_{i=1}^N y_i \\ &= \frac{0 + \dots + 1 + 0 + \dots + 1 + 0 + \dots + 1}{27} \\ &= \frac{3}{27}\end{aligned}$$

$$\begin{aligned}S^2 &= \frac{\sum_{i=1}^N (y_i - \bar{y})^2}{N - 1} \\ &= \frac{(0 - .11)^2 + \dots + (1 - .11)^2}{27 - 1} \\ &= \left(\frac{27}{27 - 1} \right) [\bar{y} \cdot (1 - \bar{y})]\end{aligned}$$

Outline

Descriptive Regression with a Binary Covariate

Conditional Means

- ▶ In this class, we focus on a single “outcome” variable.
- ▶ All additional variables adjust how we consider the outcome, often by conditioning.

Conditional Means


- ▶ In this class, we focus on a single “outcome” variable.
- ▶ All additional variables adjust how we consider the outcome, often by conditioning.
- ▶ Typically, we consider the mean of y conditional on values of other variables (called covariates).

Conditional Means

- ▶ In this class, we focus on a single “outcome” variable.
- ▶ All additional variables adjust how we consider the outcome, often by conditioning.
- ▶ Typically, we consider the mean of y conditional on values of other variables (called covariates).
- ▶ Consider the binary covariate Party (Rep 0, Dem 1).

Justice	CLlib	Party
Black	74.2	1
Reed	33.3	1
Thomas	25.6	0


From Univariate to Regression



From Univariate to Regression

pdf



2



Binary Covariate Regression

pdf

2



Conditional Mean/Median Discussion


In pairs, discuss ...

1. What do the mean and median represent in the context of the CLib variable for the justices of each party?
2. What does it mean for the mean to be larger for the Democrats than the Republicans?
3. What does it mean for the difference in means to be larger than the difference in medians?

Scatterplot Representation



OLS Regression




Intercept and Slope Discussion

In pairs, discuss ...

1. The substantive interpretation of the intercept and slope of the line in terms of justices for each party.
2. What does it mean for the mean to be larger for the Democrats than the Republicans?
3. What does it mean for the difference in means to be larger than the difference in medians?


Conditional Standard Deviations




Special Case: Binary Outcomes

Justice	Ur	Party
Black	0	1
Reed	0	1
Thomas	1	0


Binary Outcome OLS Regression



Binary Outcome OLS Regression



Conditional Standard Deviations



Binary Outcome OLS Discussion

In pairs, discuss ...

1. The substantive interpretation of the intercept and slope of the line in terms of justices for each party.
2. Why is the SD smaller for Democrats than Republicans (hint: look at the variance formula for binary variables)

Summary